

```

In [1]: # Author: Chidura Santosh
        # Importing required Libraries

import numpy as np
import pandas as pd
import scipy.stats as stats
import matplotlib.pyplot as plt
import sklearn
import seaborn as sns
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
from sklearn.datasets import load_boston

```

```

In [2]: # Loading Data and creating data frame with the data
boston = load_boston()
features = pd.DataFrame(boston.data, columns=boston.feature_names)
targets=boston.target

```

```

In [3]: #Displaying the first 5 header records
features.head()

```

Out[3]:

	CRIM	ZN	INDUS	CHAS	NOX	RM	AGE	DIS	RAD	TAX	PTRATIO	B	LSTAT
0	0.00632	18.0	2.31	0.0	0.538	6.575	65.2	4.0900	1.0	296.0	15.3	396.90	4.98
1	0.02731	0.0	7.07	0.0	0.469	6.421	78.9	4.9671	2.0	242.0	17.8	396.90	9.14
2	0.02729	0.0	7.07	0.0	0.469	7.185	61.1	4.9671	2.0	242.0	17.8	392.83	4.03
3	0.03237	0.0	2.18	0.0	0.458	6.998	45.8	6.0622	3.0	222.0	18.7	394.63	2.94
4	0.06905	0.0	2.18	0.0	0.458	7.147	54.2	6.0622	3.0	222.0	18.7	396.90	5.33

```
In [4]: #summary of the boston data statistics
features.describe()
```

Out[4]:

	CRIM	ZN	INDUS	CHAS	NOX	RM	AGE	DIS	RAD	TAX	PTRATIO	B
count	506.000000	506.000000	506.000000	506.000000	506.000000	506.000000	506.000000	506.000000	506.000000	506.000000	506.000000	506.000000
mean	3.593761	11.363636	11.136779	0.069170	0.554695	6.284634	68.574901	3.795043	9.549407	408.237154	18.455534	356.674
std	8.596783	23.322453	6.860353	0.253994	0.115878	0.702617	28.148861	2.105710	8.707259	168.537116	2.164946	91.294
min	0.006320	0.000000	0.460000	0.000000	0.385000	3.561000	2.900000	1.129600	1.000000	187.000000	12.600000	0.320
25%	0.082045	0.000000	5.190000	0.000000	0.449000	5.885500	45.025000	2.100175	4.000000	279.000000	17.400000	375.370
50%	0.256510	0.000000	9.690000	0.000000	0.538000	6.208500	77.500000	3.207450	5.000000	330.000000	19.050000	391.440
75%	3.647423	12.500000	18.100000	0.000000	0.624000	6.623500	94.075000	5.188425	24.000000	666.000000	20.200000	396.220
max	88.976200	100.000000	27.740000	1.000000	0.871000	8.780000	100.000000	12.126500	24.000000	711.000000	22.000000	396.900

```
In [5]: # To display the information of teh data frame
features.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 506 entries, 0 to 505
Data columns (total 13 columns):
CRIM      506 non-null float64
ZN        506 non-null float64
INDUS     506 non-null float64
CHAS      506 non-null float64
NOX       506 non-null float64
RM        506 non-null float64
AGE       506 non-null float64
DIS       506 non-null float64
RAD       506 non-null float64
TAX       506 non-null float64
PTRATIO   506 non-null float64
B         506 non-null float64
LSTAT     506 non-null float64
dtypes: float64(13)
memory usage: 51.5 KB
```


```
In [6]: # To check and display the sum of null values  
features.isnull().sum()
```


```
Out[6]: CRIM      0  
        ZN        0  
        INDUS    0  
        CHAS     0  
        NOX      0  
        RM       0  
        AGE      0  
        DIS      0  
        RAD      0  
        TAX      0  
        PTRATIO  0  
        B        0  
        LSTAT    0  
dtype: int64
```

```
In [7]: # Creating heat map for data visualization with correlation and coefficients
```

```
f, ax = plt.subplots(figsize=(15, 15))  
sns.heatmap(data=features.corr(), annot=True)
```

```
Out[7]: <matplotlib.axes._subplots.AxesSubplot at 0x25dcb4c1588>
```






```
In [8]: #from the above Heat map RM and lstat have higher positive and negative corelation with price values
#RAD TAX have high multi colinearity same goes with DIS and age,dis and nox column values
# displaying all the column names
features.columns
```

```
Out[8]: Index(['CRIM', 'ZN', 'INDUS', 'CHAS', 'NOX', 'RM', 'AGE', 'DIS', 'RAD', 'TAX',
               'PTRATIO', 'B', 'LSTAT'],
              dtype='object')
```

```
In [9]: # Plotting the 'RM','PTRATIO','LSTAT' against Proce
plt.figure(figsize=(20, 5))
```

```
# iterating for each column 'RM','PTRATIO','LSTAT'
for i, col in enumerate(['RM','PTRATIO','LSTAT']):
    plt.subplot(1, 3, i+1)
    x = features[col]
    y = targets
    plt.plot(x, y, 'o')
    # Create regression Line
    plt.plot(np.unique(x), np.poly1d(np.polyfit(x, y, 1))(np.unique(x)))
    plt.title(col)
    plt.xlabel(col)
    plt.ylabel('prices')
```



```
In [10]: # Creating different sizes of room based on their sizes using Quantile-based discretization function

dd=pd.qcut(features.RM,q=[0, .25, .5, .75, 1.])
pd.Categorical(dd)
dd=pd.get_dummies(dd,prefix='RM_')
dd=dd.rename(index=str, columns={"RM__(3.56, 5.885]":"Very_Small_room", "RM__(5.885, 6.208]": "Small_room", "RM__(6.208, 6.
features.index=dd.index
features['Very_Small_room']=dd['Very_Small_room']
features['Small_room']=dd['Small_room']
features['Medium_room']=dd['Medium_room']
features['Large_room']=dd['Large_room']
```

```
In [11]: # Creating different status using Quantile-based discretization function

dd=pd.qcut(features.LSTAT,q=[0, .25, .5, .75, 1.])
pd.Categorical(dd)
dd=pd.get_dummies(dd,prefix='status_')
dd=dd.rename(index=str, columns={"status__(1.729, 6.95]":"Least_lower_Status", "status__(6.95, 11.36]": "Medium_lower_Status",
features.index=dd.index
features['Majorly_lower_Status']=dd['Majorly_lower_Status']
features['lower_Status']=dd['lower_Status']
features['Medium_lower_Status']=dd['Medium_lower_Status']
features['Least_lower_Status']=dd['Least_lower_Status']
```

```
In [12]: # Creating Independent(X) and Dependent(Target-Y) data frames
X = features[['CRIM', 'ZN', 'INDUS', 'CHAS', 'NOX', 'AGE', 'DIS', 'RAD', 'TAX',
             'PTRATIO', 'B', 'Very_Small_room', 'Small_room',
             'Medium_room', 'Large_room', 'Majorly_lower_Status', 'lower_Status',
             'Medium_lower_Status', 'Least_lower_Status']]
Y = targets
```

Random Forest :

```
In [13]: from xgboost.sklearn import XGBRegressor
from sklearn.ensemble import RandomForestRegressor
from sklearn.pipeline import make_pipeline
from sklearn.tree import DecisionTreeClassifier
from sklearn.preprocessing import Imputer, StandardScaler
from sklearn.feature_selection import SelectFromModel
from sklearn.model_selection import train_test_split, GridSearchCV, ShuffleSplit, RandomizedSearchCV
import pickle
```

```
C:\ProgramData\Anaconda3\lib\site-packages\sklearn\ensemble\weight_boosting.py:29: DeprecationWarning: numpy.core.umat
h_tests is an internal NumPy module and should not be imported. It will be removed in a future NumPy release.
    from numpy.core.umath_tests import inner1d
```

```
In [14]: X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.25, random_state=42)
```

```
In [15]: #imputing null value of each column with the mean of that column
imput = Imputer()
X_train = imput.fit_transform(X_train)
X_test = imput.fit_transform(X_test)
```

```
In [16]: #Initialization for random forest
pipe = make_pipeline(StandardScaler(),
                    RandomForestRegressor(n_estimators=500, random_state=123))

cv = ShuffleSplit(test_size=0.2, random_state=0)

param_grid = {'randomforestregressor__max_features':['sqrt', 'log2', 10],
              'randomforestregressor__max_depth':[9, 11, 13]}

grid = GridSearchCV(pipe, param_grid=param_grid, cv=cv)
```


In [17]: *#finding feature_importance for feature selection. from it we'll be able to decide threshold value*

```
model = XGBRegressor()
model.fit(X_train, y_train)
print(model.feature_importances_)

[0.19127516 0.02013423 0.02516779 0.00503356 0.06040268 0.10067114
 0.15939598 0.02181208 0.06711409 0.06040268 0.11912752 0.02013423
 0.00167785 0.01174497 0.03691275 0.03020134 0.00671141 0.02013423
 0.04194631]
```

In [18]: `selection = SelectFromModel(model, threshold=0.01, prefit=True)`
`select_X_train = selection.transform(X_train)`
`select_X_test = selection.transform(X_test)`
deviding train and test data sets
`X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.25, random_state=42)`
`grid.fit(select_X_train, y_train)` *#training*

Out[18]: GridSearchCV(cv=ShuffleSplit(n_splits=10, random_state=0, test_size=0.2, train_size=None),
error_score='raise',
estimator=Pipeline(memory=None,
steps=[('standardscaler', StandardScaler(copy=True, with_mean=True, with_std=True)), ('randomforestregressor', RandomForestRegressor(bootstrap=True, criterion='mse', max_depth=None,
max_features='auto', max_leaf_nodes=None,
min_impurity_decr...imators=500, n_jobs=1,
oob_score=False, random_state=123, verbose=0, warm_start=False))]),
fit_params=None, iid=True, n_jobs=1,
param_grid={'randomforestregressor__max_features': ['sqrt', 'log2', 10], 'randomforestregressor__max_depth': [9, 11, 13]},
pre_dispatch='2*n_jobs', refit=True, return_train_score='warn',
scoring=None, verbose=0)

In [19]: `grid.best_params_`

Out[19]: {'randomforestregressor__max_depth': 9,
'randomforestregressor__max_features': 'sqrt'}

In [20]: `Randfor_reg = pickle.dumps(grid)`

```
In [21]: Randfor_reg = pickle.loads(Randfor_reg)
print("""RandomForest regressor accuracy is {ran}""").format(ran=Randfor_reg.score(select_X_test, y_test))
```

RandomForest regressor accuracy is 0.7990931731692065

```
In [ ]:
```